



SIMON

WHITEPAPER

A distributed habitat for open weights. People with GPUs compute inference, and the network verifies they actually did the work.

Version: 1.0 · 2026-09-17

Status: architecture + protocol core; network running on two machines (test), computation verification
DESIGNED, not implemented

Concept: Karol (Szpili) · implementation and measurements: Szpon (OpenClaw) + Claude Code

Governing rule: **no evidence, no status** (DESIGNED / IMPLEMENTED / MEASURED).

1. Summary

SIMON is a **distributed habitat for open model weights**. People with GPUs compute inference, receive a token for it, and the network verifies that they really did the work — without a single server anyone can switch off.

Today a model is a file you download. A file has nowhere to live outside somebody else's API. SIMON is the place where weights **reside**: they run, they are used, they evolve, and access to them flows from the network rather than from one cloud.

The difference from Hugging Face is the entire point of the design:

	Hugging Face	SIMON
Model	a file to download	a process that runs
State	frozen at upload time	alive, used, evolving
Criterion of existence	"somebody uploaded it"	used + healthy + safe
Value	download count	real usage and quality
Metaphor	a library	a habitat

HF stores. SIMON keeps alive.

2. Problem

Three things at once:

2.1. Open weights have no infrastructure. A publicly released model sits on one server or waits to be downloaded. If the server disappears, the model disappears. There is no mechanism giving it a durable place to run.

2.2. Access to models is centralised. Even when weights are open, practical access runs through the APIs of a handful of companies. Whoever owns the API sets the price, the limits and the data-use terms.

2.3. Consumer hardware sits idle. Consumer graphics cards are idle most of the day. BOINC and Salad showed that people are willing to donate compute — but not for free, and not without verification of the work.

What is missing: a mechanism that pays for work genuinely performed and verifiable on somebody else's hardware.

3. Thesis


THINK for verified work; a burned fee as the condition of issuance.

Three elements, inseparable:

1. **Verification** — the network must be able to check that a node actually computed, rather than merely claiming to.
2. **Lottery** — the client does not pick the node, so it cannot buy "its own".
3. **Settlement** — work pays, cheating loses the stake.

Without verification the rest is pointless: a node that signs a fake result itself passes every check.
This is the point where SIMON either stands or falls.

4. Architecture



Four layers:

Layer	Role	Technology
Client	submits work, receives the result	ACP (JSON-RPC), integrates with any harness
Coordinator	assignment, settlement, registry	Rust
Node	runs the model, assembles the receipt	Rust + vLLM (Python)
Transport	announcements, dialling, exchange	libp2p (gossipsub + request-response)

Flow:

```
client broadcasts a job (gossipsub – it does not know to whom)
↓
coordinators respond → one wins (deterministic ranking)
↓
a node with a GPU runs the model and assembles a RECEIPT
↓
the client checks the signature and receipt consistency ITSELF (it does not trust the
coordinator)
↓
[VERIFIER] recomputes activations and issues a verdict ← ROLE NOT DESIGNED
```

A receipt is a signed claim, not a proof. The node submits a claim of execution together with material for probabilistic auditing. The signature alone confirms the **author of the receipt**, not the correctness of the computation.

What the client can check itself, and what it cannot. It verifies the signature, the `job_id` and model match, and that the receipt is single-use. It **cannot** establish on its own whether the node actually computed — that requires the model weights and the compute to recreate the activations. A client without a GPU either trusts a third party or has no computation verification at all. Hence a separate role:

VERIFIER — a role that is REQUIRED but NOT DESIGNED. An independent verifier recomputes the indicated activations and issues a signed verdict. Its assignment, randomisation, compensation and collusion resistance are **required but undesigned**. We do not yet write "committee of 3-7" or a quorum — those are design hypotheses, not decisions.

Architectural decision: SIMON is an **ACP agent**, not a client. The harness sees a single agent; the swarm is invisible. One adapter → it works with Claude Code, OpenClaw, dsh and any other harness that speaks ACP. (ACP is a Zed standard, Apache-2.0 — not tied to a single vendor.)

5. Verification — the heart of the project

Three layers:

5.1. TopLoc — probabilistic activation audit.

A fingerprint of selected-layer activations gives a **signal** that the declared model processed the given input. The verifier recomputes the indicated activations and compares them with the executor's commitment.

The mechanism **is not a complete cryptographic proof of the whole inference**: its guarantees depend on the controlled layers, the sampling method, numerical tolerance and the adversary model.

Status: DESIGNED at concept level; NOT IMPLEMENTED; NOT ADVERSARIALLY MEASURED.

! **Boundary:** this is not a proof in the SNARK sense. Newer analyses describe attacks that stay **outside the controlled top-k** of the final layer. TopLoc is meant to be a **signal in the protocol**, not a verdict — hence the need for random, multi-point auditing and an independent verifier role.

! **This is not our invention — it is the state of the art.** The method is published (arXiv:2501.16007, January 2025); Prime Intellect used it to verify rollouts in INTELLECT-2 (arXiv:2505.07291 — distributed 32B training on other people's hardware). **Our work is the network around it** (settlement, node selection, penalties, economics), not the method itself.

5.2. VRF (job lottery). The client does not pick the node, so it cannot buy "its own". The draw happens per session, not per turn (justified in §6).

5.3. Reputation + penalty. Stake at risk; a false result means losing it.

The honest state of affairs: today the code verifies **only the signature**, not the activations.

`activation_hash` is a plain string — **nobody recomputes it**. A node that signs a fake hash and arbitrary text passes the entire verification. Today a receipt proves "registered node X claims so", not "it computed".

The real moment of truth: Frank computes the model, Szpon verifies, and the cheater gets caught.

6. Eavesdropping — a boundary that must be stated

Two different attacks; they must not be conflated:

Attack	Does TopLoc catch it
A node pretends to be model X and proxies to Claude/GPT	YES — activation hashes will not match the weights
A farm collects prompts for distillation	NO — the node computes honestly, it merely records

The key number (ESTIMATE, not measured):

Eavesdropper's share	Draw per turn, 50 turns	One node per session
1%	39.5%	1%
5%	92.3%	5%
10%	99.5%	10%

Where the difference comes from: agent harnesses send **the entire context on every turn** (files, history, code). So drawing per turn achieves nothing — a single turn already contains everything. **The node must be drawn ONCE PER SESSION.**

Three uncomfortable conclusions:

1. **Corporations and states own the most GPUs.** A lab with 10% of the capacity collects real programming tasks — the most valuable training data there is. **And SIMON pays it in THINK for the privilege.**
2. **The MVP public network protects code WORSE than a commercial API.** An API at least has contractual data-use terms; an anonymous node has none. The slogan "don't feed the corporations" is, in this version, **untrue**.
3. **This cannot be fixed with cryptography on consumer hardware.** FHE ≈ 3 s/token; confidential GPU computation (TEE) exists only on server-grade cards — that is, on corporate hardware.

What actually works (strongest first):

1. **Private pods** — the user selects "only nodes from my list" (their own, friends', the company's). **The only real protection on consumer hardware.** For code it should be the DEFAULT.
2. **Client-side anonymisation** before sending — limitation: the logic of the code is still visible.
3. **Canaries per node** — a leak identifies the culprit and justifies slashing the stake. Works only for public leaks.
4. **An explicit label:** "public network = treat your prompt as public".

A line of enquiry investigated and closed (so nobody returns to it): "a key applied to the model" (client-side encryption of weights/activations). Tested experimentally — **it does not work:**

"You can encrypt a pipe; you cannot encrypt a computer that computes on your behalf."

A key on the weights encrypts nothing (permuting neurons does not change the function). A key on activations breaks against ReLU and LayerNorm. Even where the mask survives, **the model is public**, so the node recovers the key by dividing out the weights. There is no third path: either FHE/MPC, or TEE.

7. Economics

What is new in SIMON is inference verification, not token economics. Token economics is the best-documented minefield in crypto — other people's mistakes are written up and can be avoided.

Ninth pitfall — wash-compute (added 2026-09-17 after external audit):

The audit proposed a remedy: "emission $\leq$ a share of net burned fees". **This does not remove the mechanism — it only shifts the threshold.** A farm that needs compute anyway still runs a positive balance: it gets the work *and* a share of emission. For an AI farm, wash-compute is not fraud but **optimization** — I compute locally instead of paying someone else. The emission threshold sets the size of the bonus; it does not close the path.

The balance of a party controlling both client and executor: **profit = emission – irreversible fee – electricity – verification cost**. Part of the executor payment returns to the same pocket. If emission exceeds the irreversible costs, the self-transaction remains profitable.

At this stage we know of no permissionless mechanism that distinguishes real demand from an economically equivalent self-transaction. Therefore emission tied to work volume remains BLOCKED.

The cleanest MVP model (no per-work emission): the client pays the executor for the service; the verifier receives a share of that same fee; the protocol takes a small settlement fee; **there is no additional emission for a completed job**. THINK can be a non-transferable settlement credit, and SIMON emission is not launched before simulation demonstrates real, unrelated demand. **This removes wash-compute profitability** — you can still assign work to yourself, but with no protocol subsidy there is nothing to extract.

Conclusion: wash-compute is unsolvable at the protocol level. The protocol cannot know whether the payer and the node are the same party — unless identity data is public, which breaks prompt privacy. This is a **fourth boundary** alongside eavesdropping, alignment and activation invertibility — not a gap to close, but a boundary to name.

Precedent:

Project	What it did	Lesson
Golem (2016)	a market for compute paid in a token	supply existed, demand did not, for years
Helium	rewards for hotspots	hundreds of thousands of hotspots, pennies in revenue; location spoofing
Filecoin	issuance for disk space	Filecoin Plus → fictitious self-dealing contracts
Livepeer	verified transcoding + penalty	technically the closest to SIMON
Helium DC / Render	burn-and-mint	a ready pattern for a "compute voucher"

The failure common to all of them: issuance to suppliers grows faster than user demand. The network grows by hardware rather than by usage — issuance funds farms, not the product.

Proposed pattern (A PROPOSAL, not a decision):

```
SIMON – scarce token, TRANSFERABLE
  issuance: fixed per epoch, divided by share of PAID,
            VERIFIED work (not by token count!)
  ⇔ wallet to wallet: OK

THINK – voucher, NON-TRANSFERABLE
  created by: BURNING SIMON, at a fixed price per unit of compute
            spent on jobs
  cannot be sold → cannot be speculated on,
  nor can a miner's priority be bought
```

Why a non-transferable voucher: the miner has priority over the buyer (the cooperative principle). If THINK could be transferred, a whale would buy THINK from miners **and their priority along with it** — the market would sidestep the rule with a single transfer.

Why fixed issuance per epoch: "X THINK per N tokens" means inflation tied to NVIDIA's progress. Fixed issuance per epoch (like difficulty in BTC) gives faster cards a larger share but **does not expand supply**.

Why a fixed price per unit of compute: the user must be able to plan a budget, and the miner pays the electricity bill in local currency. Only the scarce token floats.

Law and tax — gates, not footnotes: MiCA (a transferable token offered publicly probably triggers a white paper obligation; exchange requires a CASP licence), and personal income tax in Poland (payment in cryptocurrency for a service probably makes every job a taxable event, which kills the UX). **This needs a lawyer, not guesswork.**

Order of work: economics **after** verification. As long as the cheater is not caught, issuance would reward cheaters.

8. Governance — throttling, not banning

The problem: "safe / non-pathological" is not measurable. Somebody has to decide. A central committee is the very leviathan we are running from. The market alone is no answer either — pathology can be popular.

The resolution: the network regulates by **DEGREE**, not by verdict.

Model	What happens	Reversible?
Used, healthy, safe	full allocation of hosts and compute	—
Questionable	fewer hosts / less compute	YES
Used for attacks/harm	degraded to a minimum (not zero)	YES

Four reasons for gradation:

1. **There is no central verdict** — there is an allocation function, and it runs automatically.
2. **It is reversible** — a model that fixed its problem regains resources. A ban is a sentence for life; degradation is a **state**, not a label.
3. **It matches the principle "usage decides which model lives"** — a bad model is not forbidden; it simply goes unused and receives no resources.
4. **It matches the cooperative principle** — allocation is economic.

A gap, stated openly: how does the network know a model is being used for harm? A content classifier **contradicts the prompt-privacy principle** — to detect "harm in the content" you must see

the content. **Feasible only through behavioural detection (patterns at the network level, not the content level) and client reports.**

Honestly: graduated degradation is a **hypothesis, not a proven solution**. There is no evidence that it actually cuts off pathology — a bad actor can simply use a different model. It requires simulation.

9. What is MEASURED (not estimated)

Component	Value	Source
TTFT	1.74 s	MEASURED (vLLM, prefix caching OFF)
Generation, 27B on an RTX 3090	21.8 tok/s	MEASURED
Prefill (sessions of 500→36,500 tokens)	~2500-2900 tok/s, stable to 99% of the limit	MEASURED
Error of the character heuristic (Polish prose)	<2%	MEASURED
Error of the heuristic (dense JSON)	~48%	MEASURED
Error of the heuristic (base64)	~65%	MEASURED
Error of the heuristic as a function of length	a constant percentage; independent of length	MEASURED
RAM spill above the threshold	does not occur (the KV pool is pre-allocated at startup)	MEASURED to 99% of the limit
Network (transport)	~12 tok/s	! ASSUMPTION, not measured
TopLoc verification	~1-3 s	! estimate from the paper

What the heuristic measurements tell us: the error depends on **the type of content, which is not known in advance** (prose/JSON/base64), not on length. **It cannot be patched with a single multiplier** — the only way out is to ask the tokeniser for the truth (`/tokenize`). That closed the question of whether a heuristic would do.

A 500-token answer: ~25-30 s — and the user **sees nothing for the whole duration**, because the result is shown only after verification.

! **The biggest thing to reconsider.** "Verify before showing" destroys the UX at 22 tok/s. The DeepSeek API returns the same thing in 2-3 s and streams it. The alternative: **optimistic delivery** (stream immediately, sample-check afterwards, and slash the cheater's stake). **These numbers make it worth revisiting. This is open, not settled.**

10. Boundaries — stated openly, because a defence may not stay silent

Problem	Status
A cheater faking the result	! DESIGNED (TopLoc) — the code does not verify activations , so today it is not caught
Eavesdropping at the node	✗ not fixable on consumer hardware (FHE ~3 s/token, TEE server-grade only)
A node posing as a proxy to Claude/GPT	! DESIGNED — TopLoc catches this provided verification works
Distillation farm (hoarding prompts)	✗ undetectable — it computes honestly
Invertibility of activations	✗ a prompt can be reconstructed from activations (arXiv:2505.18332) — "the node sees activations, not text" is not a privacy guarantee
Semantic alignment (is the model poisoned?)	✗ OPEN — deliberately unresolved
Sybil (one operator = 1000 nodes)	! requires simulation
Prompt privacy on the public network	✗ one node sees the text (MVP decision)

Conclusion for the user: if your prompt is valuable, use **private pods**. The MVP public network protects code **worse than a commercial API**.

11. Supply and demand

The bottleneck is DEMAND, not supply. People are happy to mine (BOINC, Salad), but why would anyone **buy** THINK when the DeepSeek API costs ≈ \$0.60 per million tokens and answers ten times faster?

Real sources of demand:

1. **Miners spend their own tokens** — the only natural demand from day one
2. Censorship resistance / no single provider
3. Privacy — **but the MVP decision (one node sees the prompt) undermines it**. It cannot be sold as "private".

The model base — card memory sets the supply (Q4 ≈ 0.6 GB per billion parameters):

Card	Model in practice
8 GB	7-8B
12 GB	≤14B
24 GB	27-32B dense / MoE ~30B
48-96 GB	70B, larger MoE
~475 GiB (DeepSeek V4.1)	out of reach — needs local pods, not a WAN

Consequence: in its first year SIMON sells open models up to ~30B, **not the frontier**. Acceptable for code, summaries and batch processing. **It does not compete with the best agentic models.**

Stages:

Stage	Active nodes	Assumption
Closed alpha	2-10	our own machines + friends
Public alpha	hundreds-thousands of installs, 50-300 active after 30 days	the usual post-launch drop-off
Token on the market	thousands quickly, but farms and speculators	sybil + law (MiCA, KYC/AML) are the gates

12. Implementation status

Code: ~5800 lines of Rust (+ tests). Protocol core, transport (libp2p), coordinator, node, harness.

Done and verified live (2026-09-17):

Element	Evidence
Context-limit gate based on truth (<code>/tokenize</code>)	smoke test: declared 1000 vs 1400 in reality → controlled refusal
Chunking + map-reduce for files above a node's limit	live: 20/20 chunks, reduction 20→2→1, zero timeouts
Agent→agent chain (CODER→TESTER)	mock + live, both stages with receipt verification
Per-process <code>peer_id</code> (previously identical for every node)	live: two nodes → two different IDs
Two concurrent agents on two GPUs	live: Szpon/qwen3.8-27b + Franko/Bielik
Transport: libp2p's default 10 s timeout (a hidden bug)	found live, raised to 300 s
Portability (macOS)	audit: no native-C dependencies; <code>temp_dir()</code> fix

What is missing:

Element	Status
Activation verification (TopLoc)	DESIGNED — this is the moment of truth
A ledger with transaction finality	BurnProof is synthetic today
The ACP layer	designed, not implemented
Simulations (sybil, economics)	not run
Model router (sleep mode)	a plan, not code

Honestly: 50-80% of today's code will be rewritten before production. That is normal, and it is the reason to freeze the format only after verification.

13. Roadmap

M0 (holes in the protocol) → **M1** (transport + binary) → **M2** (two nodes, real vLLM) → **M2.5** (TopLoc + a cheater caught) → **M2.7** (format freeze + review) → **M2.9** (10-12 cases) → **M3** (the ACP layer).

M2.5 is the moment of truth. Everything before it is preparation; everything after it is expansion.

Until a cheater is caught, we do not know whether SIMON works.

What we deliberately do NOT do: federated learning (nobody needs it for an MVP), guards per specialisation (measure first), "context before model" (a control test showed a workaround, not a fix), a coordinator registry as governance (a static file is a deliberate MVP simplification).

14. Principles that came out of this work

1. No evidence, no status. Three statuses instead of "done": - **DESIGNED** — the decision exists - **IMPLEMENTED** — only with a reference to the test that **failed before the fix** - **MEASURED** — only with a reference to the file containing the result

Reason: four times in a single session a status was written down faster than it was checked. The roadmap claimed "Computational Integrity DONE" — the code did not implement it.

2. Review before the format is committed, not after. A reviewer found 4 holes in code with 50 green tests, then 3 more (a signed `OrderAccepted`, an `order_id` derived from a counter, and the absence of computation verification).

3. Honesty about boundaries is part of the product. A document that stays silent about what it cannot do is worse than no document at all — because people make decisions based on it.

15. Open questions

1. **Should the result be shown before verification?** (UX vs verification — unresolved)
 2. **Does SIMON publicly promise "your data does not go to corporations"?** Honestly, that can be promised **only in private pods**.
 3. **Where does the signal "this model is used for harm" come from?** (behaviour / reports — never content)
 4. **How is model degradation measured?** (loss of hosts / queue priority / stake multiplier)
 5. **Does graduated degradation actually cut off pathology, or merely move it to another model?** (requires simulation)
 6. **What is the minted:burned ratio in issuance?** (requires simulation)
 7. **Does a public network make product sense** at all, given that it protects worse than an API? (private pods as the answer)
-

16. In one sentence

SIMON is the place where open weights reside — people with GPUs compute, the network checks that they really did, and a model lives for as long as it is used, healthy and safe.

Decision record: `memory/decisions/` (9 architectural files + 70+ D-* entries). *Roadmap:* `ROADMAP.md`.
Code: `crates/` (Rust, ~5800 lines). *Governing rule:* no evidence, no status.